

# IAM, Billing & Cost, Security

Thao Huy Vu

# Maharishi International University - Fairfield, Iowa



All rights reserved. No part of this slide presentation may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying or recording, or by any information storage and retrieval system, without permission in writing from Maharishi International University (MIU).

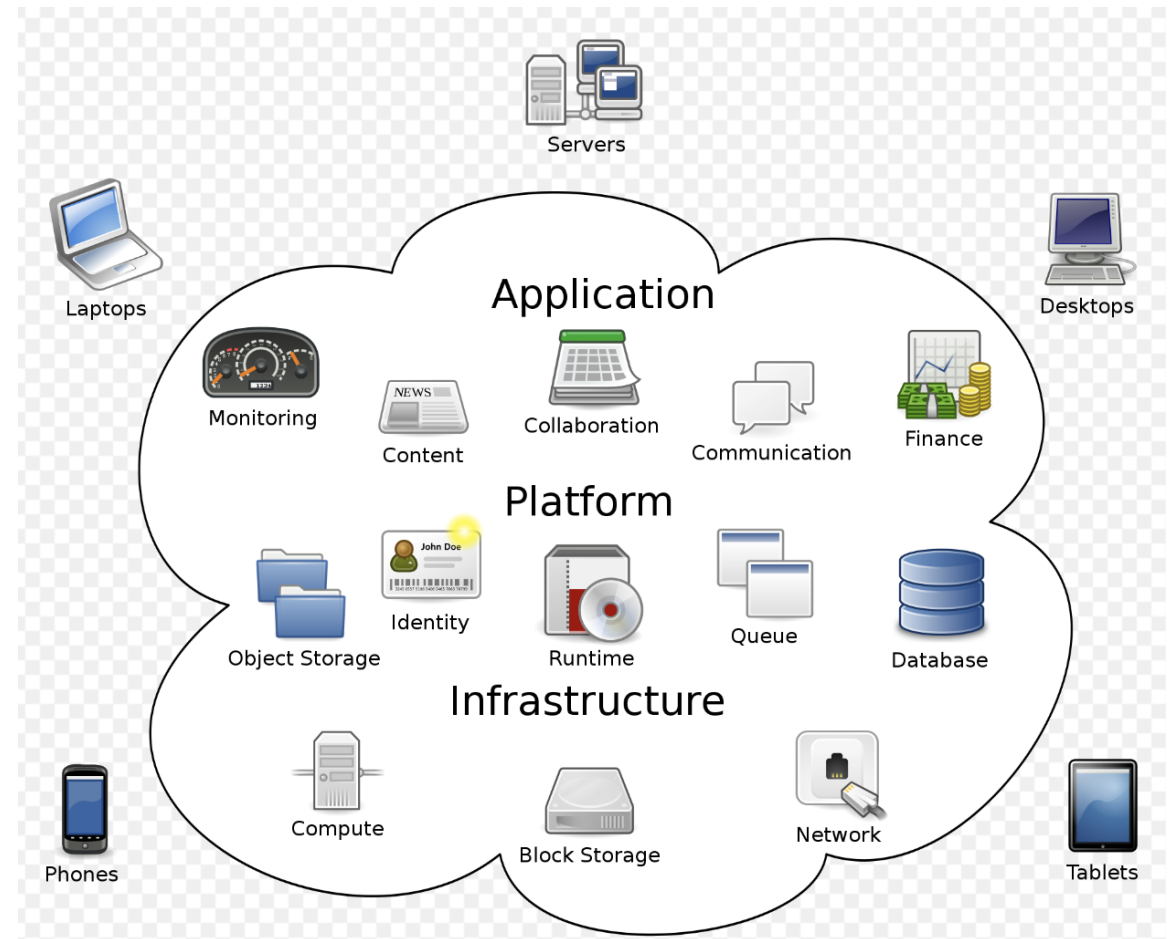
# Agenda

- AWS Overview
- IAM
- Billing & Cost
- Security

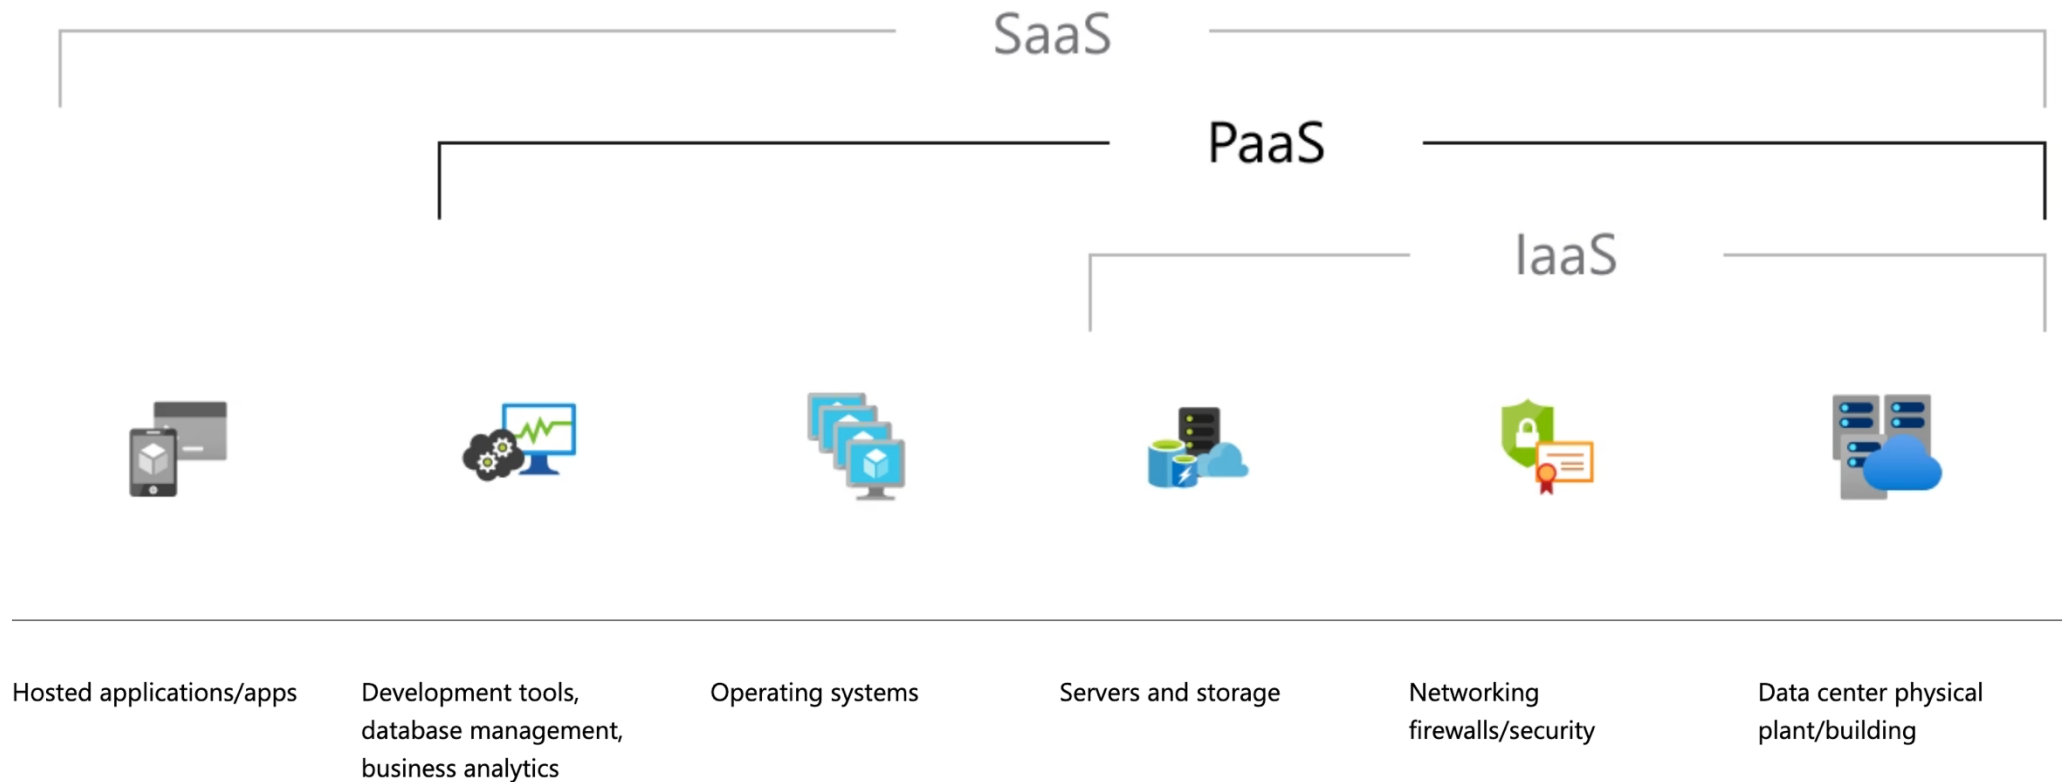
# Cloud Computing

- The group of networked elements providing services does not need to be addressed and managed by users.

[https://en.wikipedia.org/wiki/Cloud\\_computing](https://en.wikipedia.org/wiki/Cloud_computing)



# Cloud Computing model



<https://azure.microsoft.com/en-us/resources/cloud-computing-dictionary/what-is-paas/>

# Cloud Computing

- Security
- Scalability and flexibility
- Cost Efficiency
- Ease of Access
- Disaster Recovery and Data Loss Prevention
- Automatic Software updates

# Cloud Computing

- Major players
  - **AWS**
  - Microsoft Azure
  - Google Cloud Platform (GCP)
  - IBM Cloud
  - Oracle Cloud
  - Salesforce

# AWS

- Amazon Web Service.
- The largest Cloud Computing provider.
- Provide a variety of scalable computing resources and services that can be used to build, deploy, manage applications on a global network.



# Available Zones (AZ) & regions

- A region is a geographically distinct area, completely isolated.
- AZs are datacenters within a region.
- Each AZ is independent of other AZs

# AWS Global Map

**36 launched Regions**  
each with multiple Availability Zones

**114 Availability Zones**

**700+ CloudFront POPs**  
and 13 Regional edge caches


## AWS Global Infrastructure Map

The AWS Cloud spans 114 Availability Zones within 36 geographic regions, with announced plans for 12 more Availability Zones and four more AWS Regions in New Zealand, the Kingdom of Saudi Arabia, Taiwan, and the AWS European Sovereign Cloud.






## Compute

EC2  
Lightsail   
Lambda  
Batch  
Elastic Beanstalk  
Serverless Application Repository  
AWS Outposts  
EC2 Image Builder  
AWS App Runner



## Customer Enablement

AWS IQ   
Support  
Managed Services  
Activate for Startups



## Robotics

AWS RoboMaker



## Blockchain

Amazon Managed Blockchain



## Satellite

Ground Station



## Quantum Technologies

Amazon Braket



## Management & Governance

AWS Organizations  
CloudWatch  
AWS Auto Scaling  
CloudFormation  
CloudTrail



## Machine Learning

Amazon SageMaker  
Amazon Augmented AI  
Amazon CodeGuru  
Amazon DevOps Guru  
Amazon Comprehend  
Amazon Forecast  
Amazon Fraud Detector  
Amazon Kendra  
Amazon Lex  
Amazon Personalize  
Amazon Polly  
Amazon Rekognition  
Amazon Textract  
Amazon Transcribe  
Amazon Translate  
AWS DeepComposer  
AWS DeepLens  
AWS DeepRacer  
AWS Panorama  
Amazon Monitron  
Amazon HealthLake  
Amazon Lookout for Vision  
Amazon Lookout for Equipment  
Amazon Lookout for Metrics



## AWS Cost Management

AWS Cost Explorer  
AWS Budgets  
AWS Marketplace Subscriptions  
AWS Application Cost Profiler



## Front-end Web & Mobile

AWS Amplify  
Mobile Hub  
AWS AppSync  
Device Farm  
Amazon Location Service



## AR & VR

Amazon Sumerian



## Application Integration

Step Functions  
Amazon AppFlow  
Amazon EventBridge  
Amazon MQ  
Simple Notification Service  
Simple Queue Service  
SWF



## Containers

Elastic Container Registry  
Elastic Container Service  
Elastic Kubernetes Service  
Red Hat OpenShift Service on AWS



## Storage

S3  
EFS  
FSx  
S3 Glacier  
Storage Gateway  
AWS Backup

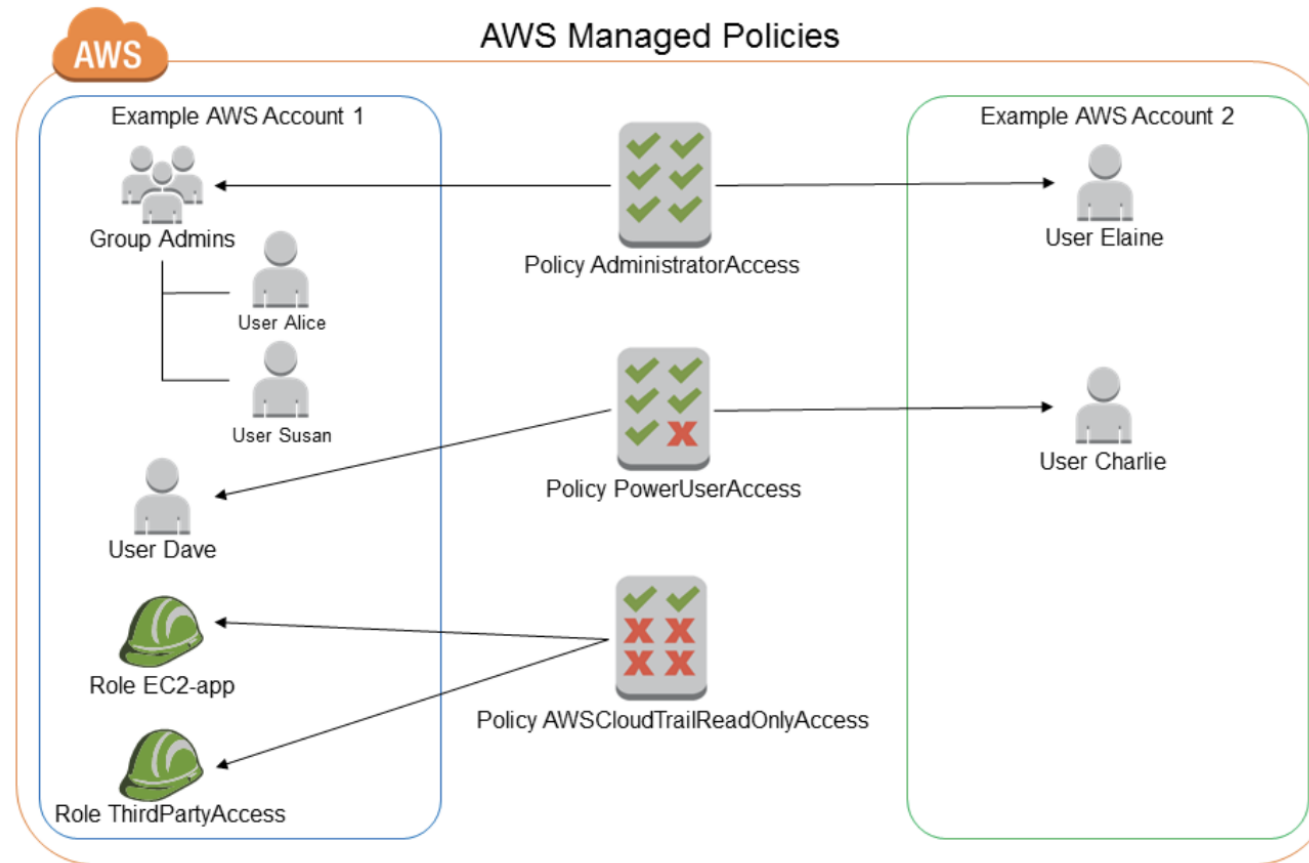
# IAM - Introduction

- AWS Identity and Access Management (IAM)
- Manage access to AWS services and resources **securely**
- Create and manage AWS users
- Allow/deny access to AWS resources
- Ensure only **authenticated** and **authorized** users can access resources

# IAM – Core components

- Users
- Group
- Role
- Policy

# IAM - Examples



# Users and Groups

- Users

- An entity represents a **person** or **service** that interacts with AWS resources.
- Provide individual **credentials** for accessing AWS services, ensuring accountability and traceability.
- Each has a name, credentials (passwords, access keys), and permissions assigned through policies.

- Groups

- A **collection** of users with specific permissions.
- Manage permissions at a group level rather than at the individual user level.
- Groups have policies attached to them which apply to all users in the group.

# Policy

- A policy is a JSON document that defines permissions to allow or deny actions on AWS resources.
- Control Access: Ensure users and roles have the appropriate access to AWS resources.
- Granular Control: Specify exactly what actions are allowed or denied.



# Policy

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "S3ReadOnlyAccess",
      "Effect": "Allow",
      "Action": [
        "s3:GetObject",
        "s3:ListBucket"
      ],
      "Resource": [
        "arn:aws:s3:::your-bucket-name",
        "arn:aws:s3:::your-bucket-name/*"
      ]
    }
  ]
}
```

# Policy

- **Identity-based** policies – Attach policies to IAM identities (users, groups, or roles)
  - Managed Policies:
    - AWS Managed Policies: Predefined policies created and managed by AWS
    - Customer Managed Policies: Created and managed by users.
    - Shared
  - Inline policies:
    - Attach directly to a single identity and belong to this identity.
    - Not shared.
- **Resource-based** policies:
  - Attach directly to a single resource(such as S3, SQS, SNS) and belong to this resource.
  - Have a **principal** element that defines who can the actions defined.


# IAM roles

- An **IAM Role** is a set of **permissions** in AWS that **any trusted entity** (like a user, EC2 instance, Lambda function, or even another AWS account) can **assume temporarily** to perform actions.
- Components:
  - Trust Policy: Defines who is allowed to assume the role.
  - Permission Policy: Defines what actions and resources the role has access to.
- Common Use Cases:
  - EC2 Instance Roles: Allow EC2 instances to access AWS services.
  - Cross-Account Access: Allow users from one AWS account to access resources in another account.
  - Service Roles: Allow AWS services like Lambda or ECS to access other AWS services.

# IAM role

An IAM user and role have one thing in common, permission policies.

When the service in the trust policy assumes the role, AWS STS (Security Token Service) returns **temporary** tokens (access key id, secret access key, and **session token**), and those tokens are rotated automatically.



# Billing & Cost

- **AWS Billing and Cost** refer to how AWS tracks, charges, and manages expenses for the cloud services you use.
- **Billing Dashboard** — View invoices and payments.
- **Cost Explorer** — Analyze and forecast costs.
- **AWS Budgets** — Set spending limits and receive alerts.
- **Cost Anomaly Detection** — Get notified about unexpected cost spikes.

# Billing Alerts (CloudWatch Alarms for Billing)

- Set up using **CloudWatch**.
- Alerts you when your total AWS charges exceed a specified threshold.
- Works on an **account-wide** level.
- Typically used for basic cost monitoring.
- Limited in flexibility (e.g., can't track specific service or linked account costs).

# AWS Budgets

- More advanced and customizable.
- Allows you to set budgets for **specific services, linked accounts, regions, or tags**.
- Can track actual costs, **forecasted costs**, and usage.
- Supports different types of budgets:
  - **Cost Budget** (total spend)
  - **Usage Budget** (specific resource usage)
  - **Reservation Budget** (Reserved Instances and Savings Plans utilization)
- Can send alerts when costs approach or exceed the budget.
- Supports multiple recipients and integration with AWS Chatbot, SNS, or emails.

# Billing and Cost

- Set **Billing Alerts**
- Set **Budget Alerts**.
- Use **Cost Allocation Tags**.
- Review costs regularly in **Cost Explorer**.
- Right-size your resources to avoid over-provisioning.



# Regular AWS accounts

- Homepage: <https://aws.amazon.com/>
- Full access to all services without limitation.
- Long-term access
- Production-level project
- Practical experience: Gain industry-like exposure.
- Credit/Debit card required.
- Free tier available: Up to 12 months for new accounts.
- **Caution: If your usage exceeds the free tier limits, you may incur unexpected charges reflected in your credit card. You are solely responsible for this charge.**

# AWS Free Tier

- **Amazon EC2:**
  - 750 hours per month of Linux or Windows t2.micro/t3.micro instances (for the first 12 months).
  - 750 hours of public IPv4 address usage per month.
- **Amazon S3:**
  - 5 GB of storage in the S3 Standard storage tier (for the first 12 months).
  - 2000 PUT, POST, COPY, or LIST requests per month.
  - 20,000 GET requests per month.
  - 15 GB of data transfer out per month.
- **Other Services:**
  - **AWS Lambda:** 1 million free requests per month.
  - **Amazon RDS:** 750 hours of a t2.micro instance.
  - **Amazon Simple Email Service (SES):** Up to 3,000 message charges free each month for the first 12 months.
  - **Amazon DynamoDB:** First 25 GB of storage and first 10 custom CloudWatch metrics are free.
  - **Amazon OpenSearch Service:** 750 hours per month of a single-AZ t2.small.search or t3.small.search instance.

# AWS Account

- Best practices to ensure cost-efficiency
  - Frequently monitor your resource usage.
  - Setup billing alert.
  - Check the bill daily.
  - Shutdown all unused instances/resources:
    - EC2 (Elastic Compute Cloud).
    - RDS (Relational Database Service), Aurora
    - ECS/EKS (Elastic Container Service/Elastic Kubernetes Service).
    - EBS (Elastic Block Store).
    - EFS (Elastic File System).
    - Unused Elastic IP addresses (EIP)
    - DocumentDB
    - DynamoDB.
    - CloudFront.
    - AWS Lambda.
    - SNS/SQS
    - VPC: NAT Gateway, Data transfer, VPN connections. You can keep the default VPC, which is already connected to the internet, so there is no need to create a new one when launching EC2 or other instances. Additionally, the default VPC does not incur significant cost.
    - S3: Particularly important when storing vast amounts of data.

# Best practices to secure AWS account and apps

1. Enable Multi-Factor Authentication (MFA) for users.
2. Use a user for deployment.
3. Never use the root account for deployment.
4. Disable the access keys of the root account.
5. Least Privilege Principle

# Reference

- AWS: <https://docs.aws.amazon.com>
- ChatGPT: <https://chatgpt.com>
- Google AI: <https://gemini.google.com>
- Practical Tutorials: <https://thaovu.org>